



MariaDB vs Neo4j

Comparative Performance Analysis of Relational and NoSQL Database Systems

Cristian Santaroni 1800659

Francesco Fortunato 1848527

Data Management Course

Master in Engineering in Computer Science

A.y. 2022/2023



OVERVIEW

Topic

Pick up a relational database system, a noSQL database system and a dataset, to compare their performances w.r.t. the execution of some relevant queries.

Goal

Find queries where relational database systems perform better and queries where noSQL database system perform better.



TOOLS

- Relational Database System: MariaDB
- NoSQL Database System: Neo4j
- Python
- Dbeaver





DATASET DESCRIPTION

Dataset Details:

- Clean, Structured, and Updated Football Data
- Sourced from Transfermarkt ([Kaggle link](#))

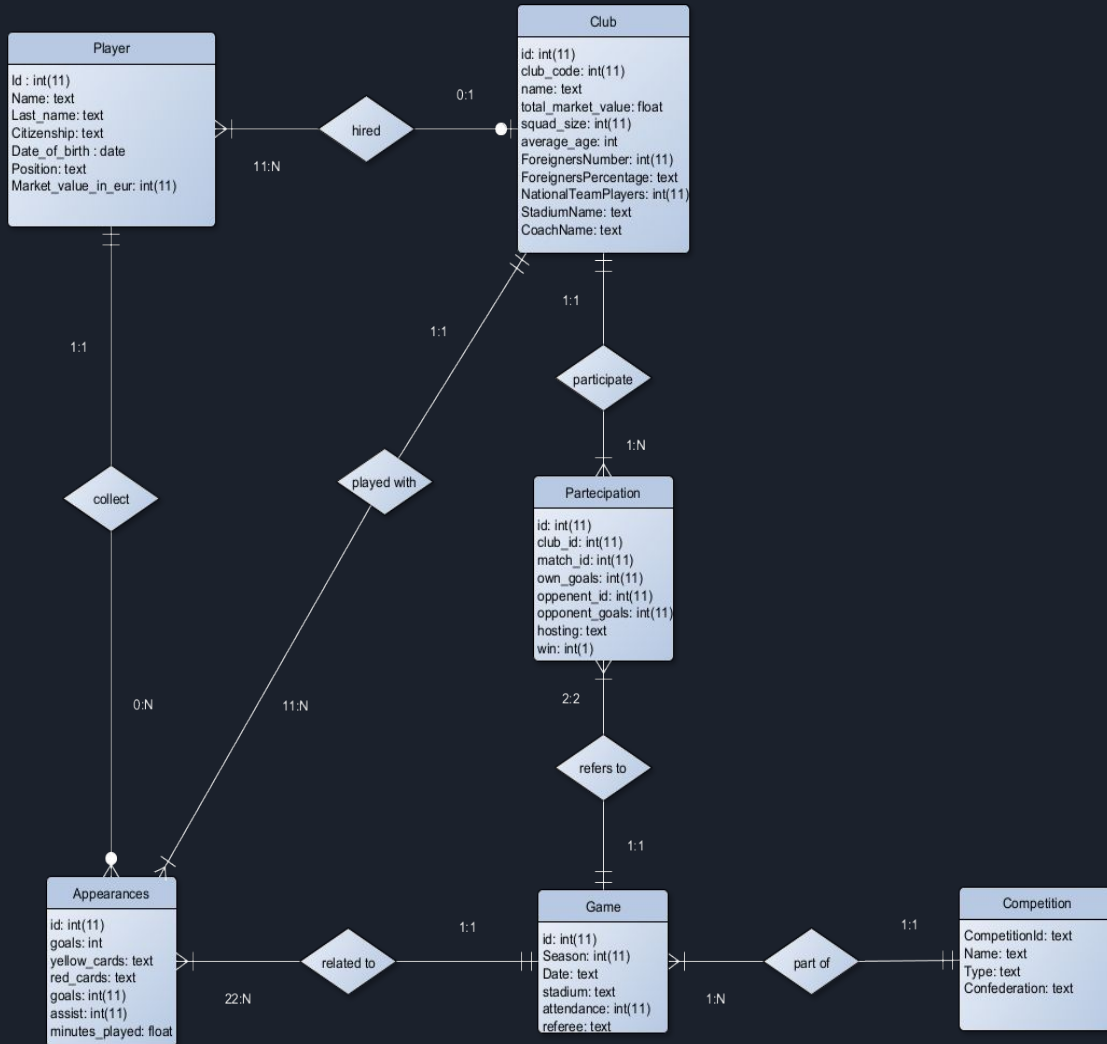
Dataset Contents:

- 63.281 Games Spanning Multiple Seasons
- 106.074 Participations Related to Games
- Representing All 43 Major Competitions
- Data on 426 Football Clubs
- Information on 29.455 Football Players
- Extensive Data on 1.179.060 Player Appearances
- Total numbers of records: 1.378.339

ER SCHEMA

Tables:

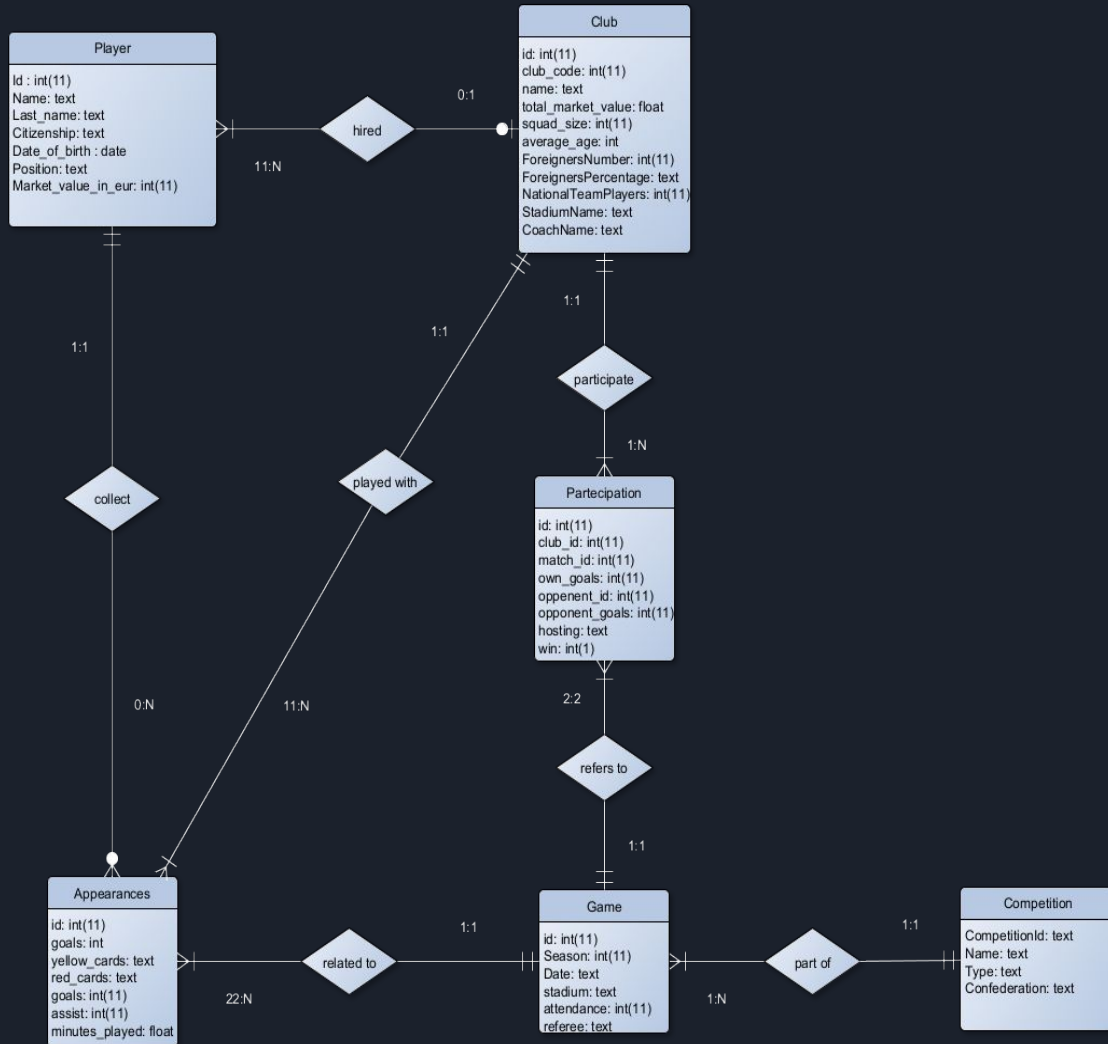
- **Player:** This table stores information about football players.
- **Appearances:** This table keeps records of player appearances in various games.
- **Game:** This table contains details about football games.
- **Club:** This table holds information about football clubs.
- **Participation:** This table links players and clubs to the games in which they participated.
- **Competition:** This table contains data about football competitions.



ER SCHEMA

Relations:

- **Collect:** player collects an appearance
- **Hired:** a player can be hired by a club
- **Refers to:** a game must have 2 clubs
- **Part of:** a game belongs to a competition
- **Played with:** an appearance is always played with a club
- **Participate:** a participation is always related to a club



DBeaver: FOREIGN KEYS

Navigatore Database X Progetti

Enter a part of object name here

localhost - localhost:3306

- Databases
 - sys
 - transformermark
 - Tables
 - appearances 220M
 - clubs 192K
 - competitions 64K
 - games 19M
 - Foreign Keys
 - Constraints
 - References
 - Triggers
 - Indexes
 - Partitions
 - participation 20M
 - players 14M

Nome tabella: games

Engine: InnoDB

Auto Increment: 0

Charset: latin1

Collation: latin1_bin

Description:

Columns	Nome	Column	Proprietario	Tabella Ref	Tipo	Oggetto Ref	On Delete	On Update
Foreign Keys	games_FK	competition_id	games	competitions	FOREIGN KEY	PRIMARY competition_id	Restrict	Restrict

Navigatore Database X Progetti

Enter a part of object name here

localhost - localhost:3306

- Databases
 - sys
 - transformermark
 - Tables
 - appearances 220M
 - clubs 192K
 - competitions 64K
 - games 19M
 - Foreign Keys
 - Constraints
 - References
 - Triggers
 - Indexes
 - Partitions
 - participation 20M
 - players 14M

Nome tabella: appearances

Engine: InnoDB

Auto Increment: 0

Charset: latin1

Collation: latin1_bin

Description:

Columns	Nome	Column	Proprietario	Tabella Ref	Tipo	Oggetto Ref	On Delete	On Update
Foreign Keys	appearances_FK_2	player_club_id	appearances	clubs	FOREIGN KEY	PRIMARY club_id	Restrict	Restrict
Foreign Keys	appearances_FK_1	game_id	appearances	games	FOREIGN KEY	PRIMARY game_id	Restrict	Restrict
Foreign Keys	appearances_FK	player_id	appearances	players	FOREIGN KEY	PRIMARY player_id	Restrict	Restrict

Navigatore Database X Progetti

Enter a part of object name here

localhost - localhost:3306

- Databases
 - sys
 - transformermark
 - Tables
 - appearances 220M
 - clubs 192K
 - competitions 64K
 - games 19M
 - Foreign Keys
 - Constraints
 - References
 - Triggers
 - Indexes
 - Partitions
 - participation 20M
 - players 14M

Nome tabella: players

Engine: InnoDB

Auto Increment: 0

Charset: latin1

Collation: latin1_bin

Description:

Columns	Nome	Column	Proprietario	Tabella Ref	Tipo	Oggetto Ref	On Delete	On Update
Foreign Keys	players_FK	current_club_id	players	clubs	FOREIGN KEY	PRIMARY club_id	Restrict	Restrict

Navigatore Database X Progetti

Enter a part of object name here

localhost - localhost:3306

- Databases
 - sys
 - transformermark
 - Tables
 - appearances 220M
 - clubs 192K
 - competitions 64K
 - games 19M
 - Foreign Keys
 - Constraints
 - References
 - Triggers
 - Indexes
 - Partitions
 - participation 20M
 - players 14M

Nome tabella: participation

Engine: InnoDB

Auto Increment: 126563

Charset: latin1

Collation: latin1_bin

Description:

Columns	Nome	Column	Proprietario	Tabella Ref	Tipo	Oggetto Ref	On Delete	On Update
Foreign Keys	participation_FK_1	club_id	participation	clubs	FOREIGN KEY	PRIMARY club_id	Restrict	Restrict
Foreign Keys	participation_FK	game_id	participation	games	FOREIGN KEY	PRIMARY game_id	Restrict	Restrict



Python scripts: importing data

In order to import all the data in our graph db, to create nodes and edges, we designed a python script. To improve the creation of the nodes and the respective edges we used a “**Uniqueness Constraint**” which ensure the uniqueness of the nodes based on the ID. This is used by Neo4j which automatically define an index to speed up the loading process and queries.

The python code can be divided in 3 parts:

- **Data Extraction** : opening all the csv files taken from kaggle and organizing them into dictionaries.
- **Node Creation**: using cypher queries, nodes are created for each data type.
- **Edges Creation**: using cypher, the relationships are created between the nodes

After this phase, data consistency was achieved with a total number of nodes of 1.378.339 and 3.842.064 relationships

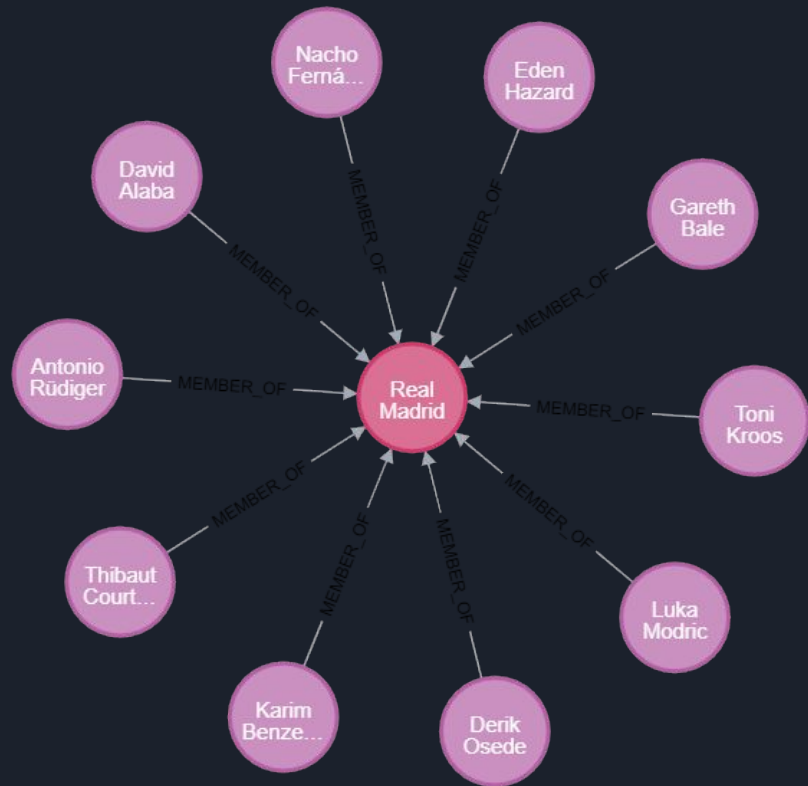
Neo4j

Node labels

* (11) Player (10) Club (1)

Relationship types

* (10) MEMBER_OF (10)



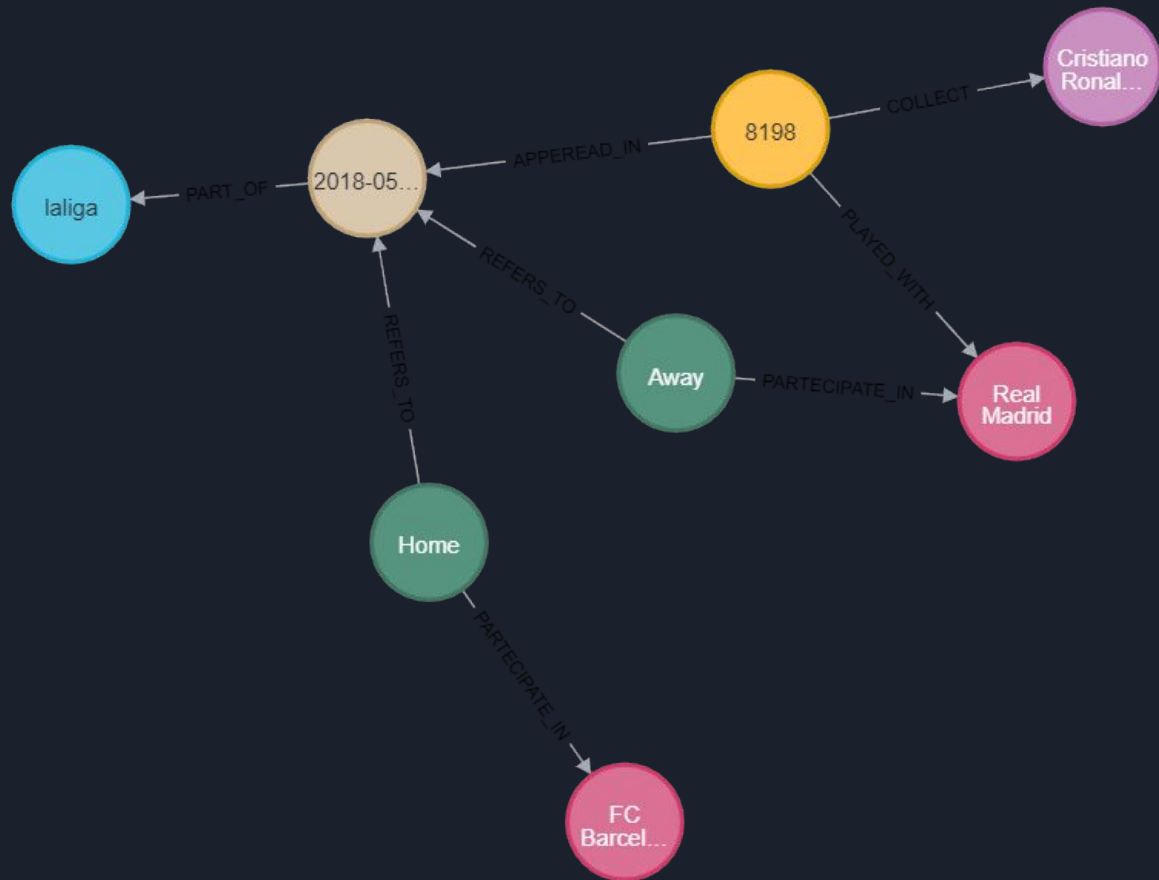
Neo4j

Node labels

* (8) Club (2) Competition (1) Player (1)
Game (1) Appearance (1) Participation (2)

Relationship types

* (8) PARTECIPATE_IN (2)
PLAYED_WITH (1) PART_OF (1)
COLLECT (1) APPEREAD_IN (1)
REFERS_TO (2)

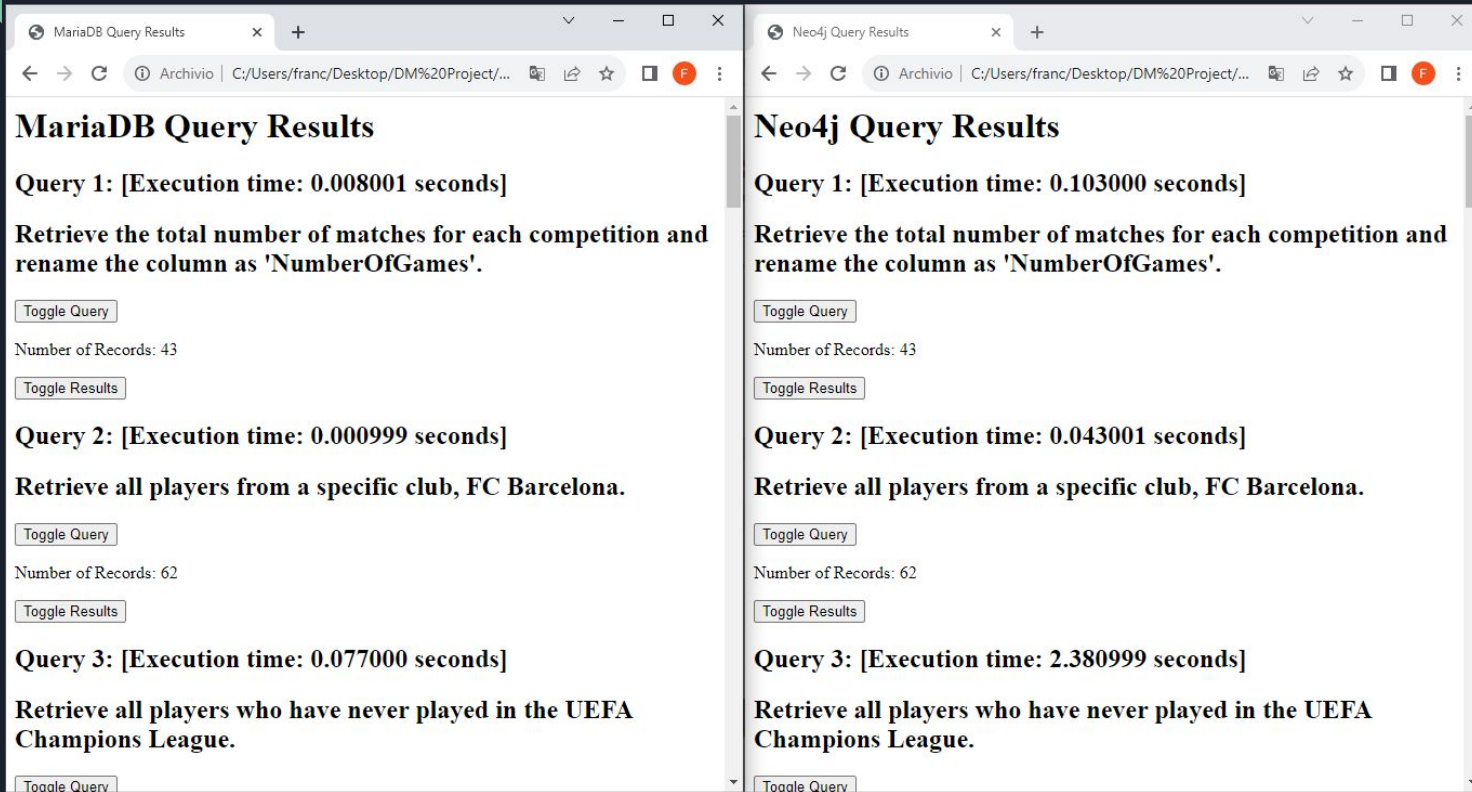




Python scripts: average execution time and demo

- For the live demo, we designed a Python script which runs 20 queries and stores the results inside two different html files, one for MariaDB and one for Neo4j. The code can be reachable [here](#).
- Regarding the performance evaluation, to measure and compare the average execution time between Mariadb and Neo4j across various queries, we've developed a dedicated Python script. The script runs each query 1000 times, and then computes the average execution time of each query. In the subsequent slides, we will present the results of this performance evaluation, providing the average execution times for each query comparing MariaDB and Neo4j.

Python scripts: Live Demo



MariaDB Query Results

Query 1: [Execution time: 0.008001 seconds]

Retrieve the total number of matches for each competition and rename the column as 'NumberOfGames'.

Number of Records: 43

Query 2: [Execution time: 0.000999 seconds]

Retrieve all players from a specific club, FC Barcelona.

Number of Records: 62

Query 3: [Execution time: 0.077000 seconds]

Retrieve all players who have never played in the UEFA Champions League.

Neo4j Query Results

Query 1: [Execution time: 0.103000 seconds]

Retrieve the total number of matches for each competition and rename the column as 'NumberOfGames'.

Number of Records: 43

Query 2: [Execution time: 0.043001 seconds]

Retrieve all players from a specific club, FC Barcelona.

Number of Records: 62

Query 3: [Execution time: 2.380999 seconds]

Retrieve all players who have never played in the UEFA Champions League.



Python scripts: Average execution time

MariaDB

```
counter = 0

for query in arr_string:
    total_time = 0
    for i in range(1000):
        start = time.time()
        cur.execute(query)
        conn.commit()
        result = list(cur)
        total_time += time.time() - start
        num_records = len(result) # Count the number of records
    counter += 1
    avg_time = total_time/1000
    print(f"Query {counter} completed. Average execution time: {avg_time:.6f} seconds")
conn.close()
```

Neo4j

```
# Function to execute query and measure execution time
def execute_query(query):
    with driver.session() as session:
        total_time = 0
        for i in range(1000):
            with session.begin_transaction() as tx:
                start_time = time.time()
                result = tx.run(query)
                # Fetch and store the results in a list
                results_list = list(result)
                tx.commit()
                total_time += time.time() - start_time
            num_records = len(results_list) # Count the number of records

        return total_time, results_list, num_records

for i, query in enumerate(queries):
    execution_time, result, num_records = execute_query(query)
    avg_time = execution_time/1000
    print(f"Query {i+1} completed. average execution time: {execution_time:.6f} seconds")
```

Average execution time (1/2)



Specs:

- AMD Ryzen 5800x3d
- Ram 16 GB
- GTX 1080

# Query	MariaDB	Neo4j
Query 1	8 ms	14.8 ms
Query 2	0.3 ms	18 ms
Query 3	80.9 ms	1.912253 s
Query 4	184.1 ms	637.721 ms
Query 5	1.604314 s	1.101904 s
Query 6	13.5 ms	299.801 ms
Query 7	1.365251 s	362.8 ms
Query 8	5.790 ms	9.4 ms
Query 9	49.865 ms	18.3 ms
Query 10	2.340601 s	116.9 ms

Average execution time (2/2)

Specs:

- AMD Ryzen 5800x3d
- Ram 16 GB
- GTX 1080

# Query	MariaDB	Neo4j
Query 11	1.754104 s	2.798699 s
Query 12	23.5 ms	287.551 ms
Query 13	998.044 ms	313.567 ms
Query 14	2.5874 s	1.466901 s
Query 15	69.4 ms	67.4 ms
Query 16	3.0753 s	2.754757 s
Query 17	172.201 ms	96.407 ms
Query 18	61.505777 s	8.48902 s
Query 19	48.671087 s	12.008614 s
Query 20	9.5 ms	28.7 ms



Conclusions: Neo4j wins

Mariadb

Pros:

- Fast Data Loading
- Excelled in simpler, structured queries
- Effective filtering

Cons:

- Sometimes queries are more complex to write
- Performance tapered off with more complex query
- Query Optimization Challenges

Neo4j

Pros:

- User-Friendly Queries
- Competitive Performance for Complex Graph-related Queries
- Most of the time is faster than RDBMS
- Consistent Performance

Cons:

- Filtering Challenges (Not great at filtering out useless nodes)
- Initial Loading Complexity.

Thank
you

